

```
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
import matplotlib.pyplot as plt
import seaborn as sns
```

```
# Importing Data from the CSV file
data = pd.read_csv('Sales.csv')
print(data.head())
```

```
↗
```

	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9


```
# Checking for null values that could require cleaning
print(data.isnull().sum())
```

```
# Getting a summary of data
print(data.describe())
```

```
↗
```

	TV	Radio	Newspaper	Sales
count	200.000000	200.000000	200.000000	200.000000
mean	147.042500	23.264000	30.554000	15.130500
std	85.854236	14.846809	21.778621	5.283892
min	0.700000	0.000000	0.300000	1.600000
25%	74.375000	9.975000	12.750000	11.000000
50%	149.750000	22.900000	25.750000	16.000000
75%	218.825000	36.525000	45.100000	19.050000
max	296.400000	49.600000	114.000000	27.000000

```
# Plotting the data
sns.pairplot(data, x_vars=['TV', 'Radio', 'Newspaper'], y_vars='Sales', height=7, aspect=0.7, kind='reg')
plt.show()
```



```
X = data[['TV', 'Radio', 'Newspaper']]
y = data['Sales']
```

```
# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

Using the Linear Regression Model

```
# Create a linear regression model
model = LinearRegression()
```

```
# Train the model using the training data
model.fit(X_train, y_train)
```

```
# Print the coefficients
print("Coefficients:", model.coef_)
print("Intercept:", model.intercept_)
```

```
➦ Coefficients: [0.05450927 0.10094536 0.00433665]
Intercept: 4.714126402214127
```

```
# Predict the sales using the testing data
y_pred = model.predict(X_test)
```

```
# Display the first few predictions
print(y_pred[:5])
```


```
➦ [17.0347724 20.40974033 23.72398873 9.27278518 21.68271879]
```

```
# Calculate the mean squared error
mse = mean_squared_error(y_test, y_pred)
print("Mean Squared Error:", mse)
```

```
# Calculate the coefficient of determination (R^2)
r2 = r2_score(y_test, y_pred)
print("R^2 Score:", r2)
```

```
# Plot the predicted vs actual values
plt.scatter(y_test, y_pred)
plt.xlabel("Actual Sales")
plt.ylabel("Predicted Sales")
plt.title("Actual vs Predicted Sales")
plt.show()
```

```
➦ Mean Squared Error: 1.4374328500000009
R^2 Score: 0.9534827934927883
```



Using the Random Forest Regression

```
from sklearn.ensemble import RandomForestRegressor
```

```
# Create a Random Forest Regressor model
model = RandomForestRegressor(n_estimators=100, random_state=24)
```

```
# Train the model using the training data
model.fit(X_train, y_train)
```

```
↳ RandomForestRegressor
RandomForestRegressor(random_state=24)
```

```
# Predict the sales using the testing data
y_pred = model.predict(X_test)
```

```
# Display the first few predictions
print(y_pred[:5])
```


```
↳ [17.479 22.761 19.698 6.614 23.098]
```

```
# Calculate the mean squared error
mse = mean_squared_error(y_test, y_pred)
print("Mean Squared Error:", mse)
```

```
# Calculate the coefficient of determination (R^2)
r2 = r2_score(y_test, y_pred)
print("R^2 Score:", r2)
```

```
# Plot the predicted vs actual values
plt.scatter(y_test, y_pred)
plt.xlabel("Actual Sales")
plt.ylabel("Predicted Sales")
plt.title("Actual vs Predicted Sales")
plt.show()
```

```
↳ Mean Squared Error: 1.4358602000000027
R^2 Score: 0.9535336865030694
```



Start coding or [generate](#) with AI.